Proof-governed autonomy for live revenue experiments

Alan N. Pham[1,*]

[1]AO Labs; institutional affiliation and coauthor list to be confirmed before submission

*Correspondence: Alan Pham

Autonomous agents are increasingly able to use tools, write messages, and execute workflows, but deployed claims of autonomy are difficult to evaluate when actions, constraints, market signals, and commercial outcomes are scattered across product analytics and human memory. We report Relay, an evidence-locked deployment of a notes-first revenue agent governed by an explicit proof contract. The system converts a broad commercial goal into a state machine with recorded bottlenecks, permitted actions, forbidden actions, proof deadlines, buyer-intent signals, payment outcomes, and compact research-journal entries. At the evidence lock, Relay exposed nine production routes, had all eleven checked environment requirements present, retained 30-day telemetry including 56 page views, 8 link clicks, 3 checkout clicks, and 3 lead-capture events, and recorded four compact research-journal entries. Crucially, it had not yet demonstrated autonomous profitability: gross revenue was USD 0.00, non-test payments were zero, and the active hard-paid-test experiment was 0 of 10 sends before its next business-hour window. The contribution is therefore not a claim that the system already makes money. It is a falsifiable protocol for making autonomous commercial experimentation auditable enough that future success or failure can be interpreted.

1 Introduction

Autonomous software agents are moving from demonstrations into deployed settings where they can observe, decide, and act through external tools (*1–5*). This creates a new evaluation problem. A tool-using system can appear autonomous because it sends messages, records data, or triggers payments, but a scientific claim of autonomy requires more than action. The system must preserve the state that justified action, the constraints that limited action, and the evidence that should update the next action.

This requirement is familiar in closed-loop experimentation. In autonomic computing, a controller monitors, analyses, plans, and executes under an explicit knowledge model (*6*). In closed-loop materials discovery, a system chooses the next experiment, records the result, and updates the search path (*7*). Commercial agent deployments need a related discipline, but with different risks: buyer privacy, outreach ethics, payment accounting, internal-test contamination, and human operator intervention. Without an explicit proof contract, a founder can repeatedly change the market, copy, price, or volume and still describe the whole process as one autonomous system.

Relay was built around a deliberately narrow offer: a buyer sends messy sales-call notes and receives a client-ready follow-up packet. The user-facing product is simple, but the research object is broader. Relay connects a product surface, buyer-intent telemetry, payment links, intake capture, outbound acquisition records, production-event routes, a success governor, an autonomous money loop, and a research-journal endpoint. The system is designed to keep the operator out of routine replanning during a proof window while preserving enough evidence to reconstruct what happened.

Here we frame Relay as a proof-governed autonomy system. The term denotes a closed loop in which the agent can act only under a recorded proof state: a measurable objective, a current bottleneck, a permitted next action, a forbidden-action set, a deadline, and outcome fields that define whether the action succeeded. This paper is evidence-locked to a saved production snapshot rather than to a live dashboard that can change after drafting. The resulting claim is intentionally bounded. Relay is deployed and auditable; it is not yet proven profitable.

2 Results

The Results are organised around four levels of the deployment: the proof-governance protocol, the deployed product surface, the evidence matrix at the lock, and the journaled active experiment that determines the next interpretable claim.

2.1 A proof contract for autonomous revenue experiments

Relay converts an open-ended goal—make the system generate money while the operator stays out of the loop—into a constrained experiment. The proof contract contains five parts. First, a revenue objective defines the target outcome. Second, the state layer records money, payments, queue state, active-sample progress, send capacity, exception state, and recent market signals. Third, the success governor computes the bottleneck and selects the next allowed action. Fourth, the guardrail contract blocks actions that would make the experiment uninterpretable. Fifth, the research journal records the control tick and money-loop tick.

At the evidence lock, the success state was *active_experiment_sample*. The owner mode was *out_of_loop_waiting*. The allowed action was to send 10 active leads during the approved business-hour window. The forbidden set included asking the owner to choose the next move before the proof deadline, increasing volume before buyer signal, changing more than one major variable at a time, and declaring demand failure from an execution miss. This converts an emotionally loaded commercial objective into a narrow experimental state.

Table 1. Proof contract at the evidence lock. The table separates the current state from the permitted action and the claims that remain unsupported.

Element	Recorded value	Experimental function
Objective	Reach a USD 100 weekly revenue target through paid tests	Defines the target outcome without changing the evidence.
Money state	USD 0.00 gross revenue; zero non-test payments	Prevents internal or unmatched test events from becoming a false success.
Bottleneck	Active hard-paid-test sample incomplete	Blocks premature judgement of demand.
Allowed action	Send 10 active leads inside the approved cap and window	Completes the minimum exposure before interpreting signal.
Forbidden actions	No owner-choice request, no volume increase, no multi-variable change, no failure claim from execution miss	Preserves causal interpretability.
Audit point	2026-05-04T17:00:00-04:00	Defines when the next evidence should be evaluated.

2.2 The product surface is part of the experiment

The buyer-facing surface is intentionally direct. A visitor sees the offer, submits messy notes, views a sample packet, and can start a paid one-call test. These interface elements are not only conversion elements; they are also sensors. Viewing the sample, focusing the note field, submitting notes, clicking checkout, and entering an email all create evidence used by the proof contract.

Figure 1 shows the production product surface used by the experiment. The desktop page presents the offer, output promise, sample packet preview, and note-intake form in one view. The mobile view compresses the same path into a single-column flow. The preview is included here because future claims about buyer behaviour depend on what buyers actually saw.

The returned artifact is a follow-up packet. Figure 2 shows the sample preview used on the product page. This matters because the experimental offer is not a vague artificial-intelligence claim; it is a concrete promise that rough sales-call notes can be converted into a recap, follow-up draft, open questions, and CRM-ready update.

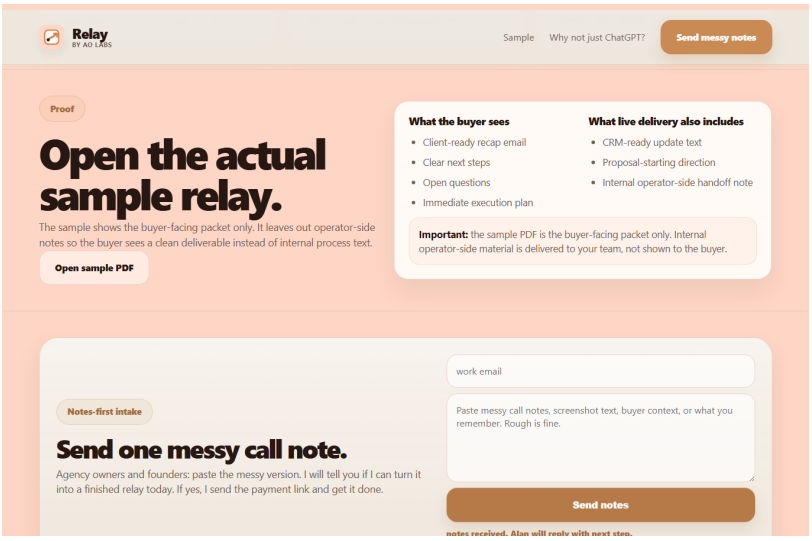

(a) Desktop product page.

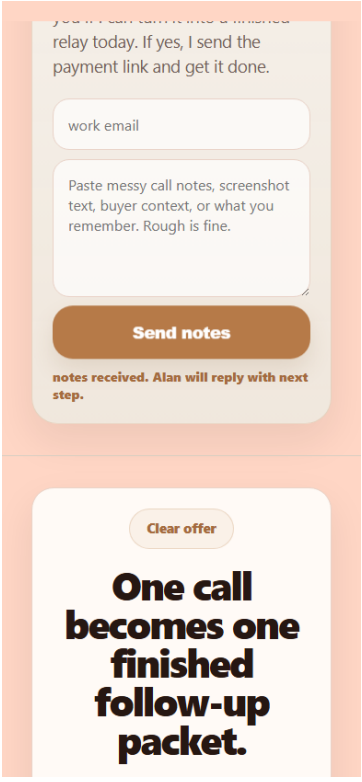

(b) Mobile product page.

Figure 1. Buyer-facing Relay surface used by the evidence-locked deployment. The surface is both a product page and an experimental instrument: note-intake focus, sample clicks, lead capture, and checkout intent are recorded as buyer-intent signals.

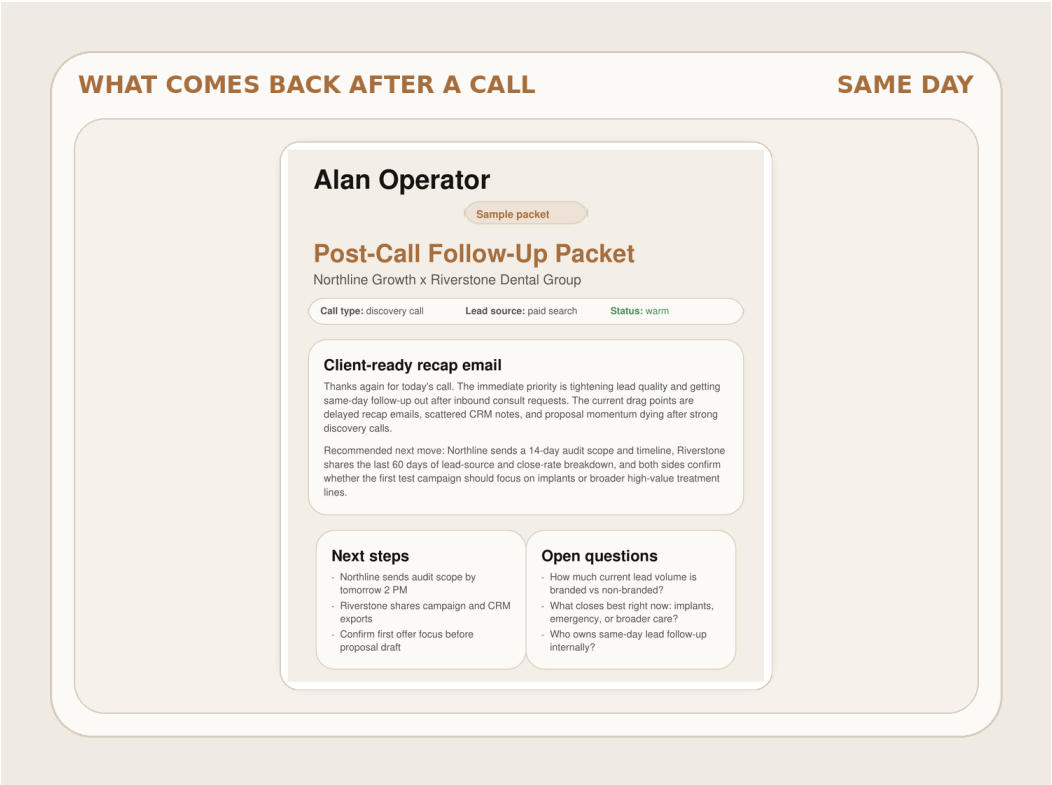

Figure 2. Relay sample-packet preview shown to buyers. The sample anchors the offer in a specific deliverable rather than in a generic automation claim.

2.3 Evidence at the lock supports deployment, not profit

The saved operational snapshot showed a live route surface and a configured backend, but the money state was still zero. Table 2 summarises the supported and unsupported claims. This distinction is the main scientific discipline of the manuscript. A stronger paper does not inflate early signals; it records where the proof stops.

Table 2. Evidence matrix at the 2026-05-04 lock. Values are descriptive and should not be read as statistical evidence of demand.

Domain	Evidence at lock	Supported claim	Unsupported claim
Route surface	Stripe and Tally webhooks, client gate, production routes, autopilot routes, intent summary, research journal	Relay was a deployed operational system, not only a static page.	Route robustness under buyer load.
Configuration	11 of 11 checked environment requirements present	Payment, intake, delivery, and tracking routes had required configuration.	Correct behaviour for every future edge case.
Intent	56 page views, 8 link clicks, 3 checkout clicks, 3 lead captures	There was measurable site interaction.	Product-market fit or statistically reliable conversion.
Filtered buyer signal	18 real page views, 1 real lead, 2 internal test leads separated	Internal activity was not treated as buyer demand.	A large external buyer sample.
Outbound	60 rolling seven-day sends, 0 replies, 0 payments, 0 send failures	Prior outreach and non-response were visible.	Outbound conversion.
Money	USD 0.00 gross revenue, zero non-test payments	Profit was not yet proven.	Autonomous money-making success.
Active experiment	0/10 active sends, 23 active due leads, 10/day cap	The next proof action was defined.	Demand failure before the sample completes.
Journal	4 compact research-journal entries	Control states were preserved for audit.	Longitudinal generalisation.

2.4 The active experiment defines the next interpretable result

At the evidence lock, the active experiment was a hard paid test. It had 0/10 progress, 10 sends remaining, 45 queued direct leads, 23 active due leads, and a remaining cap of 10. The next autonomous send window was 2026-05-04T09:30:00-04:00, and the proof audit point was 2026-05-04T17:00:00-04:00. Because the snapshot was taken before the send window, zero active sends was an expected pre-window state rather than evidence of market rejection.

The primary endpoint for the next window is non-test paid purchase count and gross revenue. Secondary endpoints are replies, checkout clicks, note submissions, lead captures, send failures, and open exceptions. If the active sample completes and no buyer signal appears, the protocol should rotate one major variable—targeting, copy, price, or volume—rather than rewriting the entire system.

If buyer signal appears, the protocol should preserve the lane and close the payment or fulfilment loop.

2.5 Compact journaling makes the system publishable as a research object

Before the journal layer, Relay could be monitored, but later reconstruction required multiple routes and human memory. The research-journal endpoint changes the status of the deployment because the system now records compact control entries. These entries are not motivational notes; they are evidence objects. They preserve the money state, active-sample progress, proof state, health state, and next action.

Table 3. Predeclared interpretation of the next active sample. The goal is not to make a dashboard feel optimistic, but to make the next result interpretable.

Observed outcome	Interpretation	Permitted next action
Paid test or strong buyer reply	The lane has buyer signal.	Fulfil, preserve the lane, and continue only within the current cap.
Checkout clicks without payment	Friction may be in price, trust, or checkout path.	Change one variable and run another small sample.
No replies or clicks after complete sample	Current lane has no signal at this sample size.	Rotate one major variable; do not change targeting, copy, price, and volume simultaneously.
Send-window failure	Execution failed before demand was tested.	Fix execution before drawing market conclusions.
Manual owner intervention	Experiment becomes contaminated.	Mark the window as interrupted and separate it from autonomous evidence.

Table 4. Latest compact research-journal entries at the evidence lock. Buyer personal data are excluded from the manuscript.

Created at	Summary
2026-05-04T04:42:05.206321	money_loop money=$0.00 payments=0 sent=0 active_delta=0/0 active=0/10 proof=not_expected next=send 10 active leads and move progress from 0/10 to 10/10
2026-05-04T04:41:57.336593	success_control money=$0.00 payments=0 bottleneck=active_experiment_sample proof=prove_active_sample health=waiting_for_proof_deadline conversion_sent=0 failures=0
2026-05-04T04:40:17.309858	money_loop money=$0.00 payments=0 sent=0 active_delta=0/0 active=0/10 proof=not_expected next=send 10 active leads and move progress from 0/10 to 10/10
2026-05-04T04:40:14.073885	success_control money=$0.00 payments=0 bottleneck=active_experiment_sample proof=prove_active_sample health=waiting_for_proof_deadline conversion_sent=0 failures=0

3 Discussion

Relay is not yet evidence that an autonomous agent can make money without an operator. It is evidence that an autonomous commercial experiment can be constrained and logged so that a later money claim would be auditable. That is the stronger contribution. The system's most important result is the discipline of refusing to count an internal test as revenue, refusing to judge demand before the active sample, and refusing to change many variables at once.

The broader implication is that agentic systems need proof contracts when deployed in domains with real-world consequences. A proof contract makes autonomy less theatrical and more scientific. It identifies what the system may do, what it must not do, when the result can be judged, and what evidence changes the next state. This structure could apply beyond Relay to other live agent deployments where market response, human intervention, or hidden tool actions make claims difficult to interpret.

The current deployment remains preliminary. It is a single system, a single offer, and a small descriptive dataset. It contains zero real revenue at the evidence lock. It does not provide inferential statistics, external replication, or a mature ethics protocol. A Nature-level version of this work would need repeated journaled rotations, non-test buyer outcomes, a comparison with manual or dashboard-only operation, a privacy-preserving dataset, independent review of code and logs, and a clear institutional position on commercial outreach and data retention.

4 Methods

4.1 Evidence lock and data sources

The manuscript is locked to three saved production snapshots: `relay_ops_snapshot_2026-05-04T0042.json`, `relay_journal_snapshot_2026-05-04T0042.json`, and `relay_intent_snapshot_2026-05-04T0042.json`. The code state used for the evidence lock was commit `be6d9104aebd10dcf42b03d54f2d0ce1b8a9043b` in repository `nalalalan/relay-app`. Raw snapshots are retained locally for audit but are not reproduced in the public manuscript because they may contain operationally sensitive fields or personal identifiers.

4.2 Revenue accounting

Revenue was read from the money-system summary and success-controller state rather than from raw acquisition status counts alone. This prevents internal or unmatched Stripe test events from being counted as real buyer revenue. At the evidence lock, the money-system summary reported USD 0.00 gross revenue and zero payments. That value is used throughout the manuscript.

4.3 Intent and internal-test filtering

Intent was reported in two forms: all recorded intent events and filtered real-event counts where available. This is necessary because early deployment testing can otherwise mimic buyer demand. The manuscript reports aggregate counts only and excludes buyer emails and session identifiers.

4.4 Autonomy policy

The autonomy policy was read from the success-controller and money-system summaries. The policy defines the current bottleneck, allowed next action, forbidden actions, proof deadline, owner-interruption state, and recovery conditions. These values are treated as experimental controls. A future protocol deviation should be recorded in the journal before the corresponding outcome is interpreted.

4.5 Statistics and sample size

No hypothesis test was performed. The available observations are descriptive: 56 page views, 18 filtered real page views, 8 link clicks, 3 checkout clicks, 3 lead-capture events, 1 filtered real lead, 60 rolling seven-day outbound sends, 0 replies, 0 payments, and 0 send failures. The active experiment target was 10 active sends, but the observed active sample at the evidence lock was 0. Future statistical claims require completed samples and a predeclared analysis plan.

4.6 AI assistance disclosure

This manuscript draft was prepared with OpenAI Codex under author direction using saved production snapshots and cited sources. The language model is not an author. The human author is responsible for factual verification, final approval, authorship decisions, and disclosure according to the target journal policy (*8*).

5 Data availability

The minimum dataset for this draft consists of the three saved JSON snapshots named in Methods plus source code needed to reproduce the route summaries and compact journal entries. Before public release, buyer emails, session identifiers, secrets, and operationally sensitive fields should be removed or transformed into privacy-preserving identifiers. A formal submission should deposit a frozen, sanitised dataset and schema in a DOI-minting repository or explain controlled-access restrictions (*9*).

6 Code availability

The implementation corresponds to repository `nalalalan/relay-app` at commit `be6d9104aebd10dcf42b03d54f2d0ce1b8a9043b`. Public release status must be confirmed before submission. If the code is central to the claims, peer reviewers should receive access to route definitions, the journal schema, a configuration template without secrets, and scripts or instructions needed to regenerate the manuscript figures from sanitised evidence.

7 Ethics and privacy

The system involves commercial outreach and buyer-intent logging. A complete study should specify lawful basis, retention period, opt-out handling, suppression-list behaviour, limits on automated follow-up, handling of personal data, and publication redaction rules. The present manuscript does not publish buyer personal data and separates internal tests from buyer signal.

8 Competing interests

The author owns or operates AO Labs and Relay and may benefit commercially if Relay succeeds. This commercial interest should be disclosed in any submission.

9 Funding

No specific funding source was recorded in the evidence used for this draft. The author should update this statement if the work falls within a grant, fellowship, institutional project, or company-funded effort.

10 Author contributions

A.N.P. conceived the system and research framing, directed the deployment goals, owns factual verification of the evidence, and is accountable for the final manuscript. AI assistance was used for drafting and formatting under author direction and is disclosed in Methods.

References

1. S. Yao *et al.*, ReAct: Synergizing reasoning and acting in language models. *International Conference on Learning Representations* (2023).
2. T. Schick *et al.*, Toolformer: Language models can teach themselves to use tools. *arXiv preprint arXiv:2302.04761* (2023).
3. J. S. Park *et al.*, Generative agents: Interactive simulacra of human behavior. *arXiv preprint arXiv:2304.03442* (2023).
4. Q. Wu *et al.*, AutoGen: Enabling next-gen LLM applications via multi-agent conversation framework. *arXiv preprint arXiv:2308.08155* (2023).
5. L. Wang *et al.*, A survey on large language model based autonomous agents. *arXiv preprint arXiv:2308.11432* (2023).
6. J. O. Kephart, D. M. Chess, The vision of autonomic computing. *Computer* **36**, 41–50 (2003).
7. A. G. Kusne *et al.*, On-the-fly closed-loop materials discovery via Bayesian active learning. *Nature Communications* **11**, 5966 (2020).
8. Nature, *Artificial intelligence policy* (2026; `https://www.nature.com/nature/editorial-policies/ai`).
9. Nature, *Reporting standards and availability of data, materials, code and protocols* (2026; `https://www.nature.com/nature/editorial-policies/reporting-standards`).